

```
In [1]: import numpy as np
import pandas as pd
```

Obtain the train and test data

```
In [2]: train = pd.read_csv('UCI_HAR_dataset/csv_files/train.csv')
test = pd.read_csv('UCI_HAR_dataset/csv_files/test.csv')
print(train.shape, test.shape)


(7352, 564) (2947, 564)
```

```
In [3]: train.head(3)
```

```
Out[3]:
```

	tBodyAccmeanX	tBodyAccmeanY	tBodyAccmeanZ	tBodyAccstdX	tBodyAccstdY	tBodyAccstdZ
0	0.288585	-0.020294	-0.132905	-0.995279	-0.983111	-0.913505
1	0.278419	-0.016411	-0.123520	-0.998245	-0.975300	-0.960300
2	0.279653	-0.019467	-0.113462	-0.995380	-0.967187	-0.978900

3 rows × 564 columns



```
In [4]: # get X_train and y_train from csv files
X_train = train.drop(['subject', 'Activity', 'ActivityName'], axis=1)
y_train = train.ActivityName
```

```
In [5]: # get X_test and y_test from test csv file
X_test = test.drop(['subject', 'Activity', 'ActivityName'], axis=1)
y_test = test.ActivityName
```

```
In [6]: print('X_train and y_train : ({},{})'.format(X_train.shape, y_train.shape))
print('X_test and y_test : ({},{})'.format(X_test.shape, y_test.shape))

X_train and y_train : ((7352, 561),(7352,))
X_test and y_test : ((2947, 561),(2947,))
```

Let's model with our data

Labels that are useful in plotting confusion matrix

```
In [7]: labels=['LAYING', 'SITTING', 'STANDING', 'WALKING', 'WALKING_DOWNSTAIRS', 'WALKING_UPSTAIRS']
```

Function to plot the confusion matrix

```
In [8]: import itertools
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
plt.rcParams["font.family"] = 'DejaVu Sans'

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=90)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```

Generic function to run any model specified

```

In [9]: from datetime import datetime
def perform_model(model, X_train, y_train, X_test, y_test, class_labels, cm_normalize=True, \
                 print_cm=True, cm_cmap=plt.cm.Greens):

    # to store results at various phases
    results = dict()

    # time at which model starts training
    train_start_time = datetime.now()
    print('training the model..')
    model.fit(X_train, y_train)
    print('Done \n \n')
    train_end_time = datetime.now()
    results['training_time'] = train_end_time - train_start_time
    print('training_time(HH:MM:SS.ms) - {}'.format(results['training_time']
    ))

    # predict test data
    print('Predicting test data')
    test_start_time = datetime.now()
    y_pred = model.predict(X_test)
    test_end_time = datetime.now()
    print('Done \n \n')
    results['testing_time'] = test_end_time - test_start_time
    print('testing_time(HH:MM:SS.ms) - {}'.format(results['testing_time']
    ))

    results['predicted'] = y_pred

    # calculate overall accuracy of the model
    accuracy = metrics.accuracy_score(y_true=y_test, y_pred=y_pred)
    # store accuracy in results
    results['accuracy'] = accuracy
    print('-----')
    print('|          Accuracy          |')
    print('-----')
    print('\n    {}'.format(accuracy))

    # confusion matrix
    cm = metrics.confusion_matrix(y_test, y_pred)
    results['confusion_matrix'] = cm
    if print_cm:
        print('-----')
        print('| Confusion Matrix |')
        print('-----')
        print('\n {}'.format(cm))

    # plot confusion matrix
    plt.figure(figsize=(8,8))
    plt.grid(b=False)
    plot_confusion_matrix(cm, classes=class_labels, normalize=True, title='Normalized confusion matrix', cmap = cm_cmap)

```

```
plt.show()

# get classification report
print('-----')
print('| Classification Report |')
print('-----')
classification_report = metrics.classification_report(y_test, y_pred)
# store report in results
results['classification_report'] = classification_report
print(classification_report)

# add the trained model to the results
results['model'] = model

return results
```

Method to print the gridsearch Attributes

```

In [10]: def print_grid_search_attributes(model):
    # Estimator that gave highest score among all the estimators formed in GridSearch
    print('-----')
    print('|           Best Estimator           |')
    print('-----')
    print('\n\t{}\n'.format(model.best_estimator_))

    # parameters that gave best results while performing grid search
    print('-----')
    print('|           Best parameters           |')
    print('-----')
    print('\tParameters of best estimator : \n\n\t{}\n'.format(model.best_params_))

    # number of cross validation splits
    print('-----')
    print('|   No of CrossValidation sets   |')
    print('-----')
    print('\n\tTotal numbere of cross validation sets: {}\n'.format(model.n_splits_))

    # Average cross validated score of the best estimator, from the Grid Search
    print('-----')
    print('|           Best Score           |')
    print('-----')
    print('\n\tAverage Cross Validate scores of best estimator : \n\n\t{}\n'.format(model.best_score_))

```

1. Logistic Regression with Grid Search

```

In [11]: from sklearn import linear_model
    from sklearn import metrics

    from sklearn.model_selection import GridSearchCV

```

```
In [12]: # start Grid search
parameters = {'C':[0.01, 0.1, 1, 10, 20, 30], 'penalty':['l2','l1']}
log_reg = linear_model.LogisticRegression()
log_reg_grid = GridSearchCV(log_reg, param_grid=parameters, cv=3, verbose=1, n
_jobs=-1)
log_reg_grid_results = perform_model(log_reg_grid, X_train, y_train, X_test,
y_test, class_labels=labels)
```

training the model..

Fitting 3 folds for each of 12 candidates, totalling 36 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 36 out of 36 | elapsed: 1.7min finished
C:\Users\Raftaar Singh\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22.
Specify a solver to silence this warning.
  FutureWarning)
C:\Users\Raftaar Singh\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:460: FutureWarning: Default multi_class will be changed to 'auto' in
0.22. Specify the multi_class option to silence this warning.
  "this warning.", FutureWarning)
```

Done

training_time(HH:MM:SS.ms) - 0:01:55.292116

Predicting test data


Done

testing time(HH:MM:SS.ms) - 0:00:00.113713

```
-----
|      Accuracy      |
|-----|
| 0.9626739056667798 |
```

```
-----
| Confusion Matrix |
|-----|


[[537  0  0  0  0  0]
 [ 1428 58  0  0  4]
 [  0 12 519  1  0  0]
 [  0  0  0 495  1  0]
 [  0  0  0  3 409  8]
 [  0  0  0 22  0 449]]
```



Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.97	0.87	0.92	491
STANDING	0.90	0.98	0.94	532
WALKING	0.95	1.00	0.97	496
WALKING_DOWNSTAIRS	1.00	0.97	0.99	420
WALKING_UPSTAIRS	0.97	0.95	0.96	471
micro avg	0.96	0.96	0.96	2947
macro avg	0.97	0.96	0.96	2947
weighted avg	0.96	0.96	0.96	2947


```
In [13]: plt.figure(figsize=(8,8))
plt.grid(b=False)
plot_confusion_matrix(log_reg_grid_results['confusion_matrix'], classes=labels
, cmap=plt.cm.Greens, )
plt.show()
```



```
In [14]: # observe the attributes of the model
print_grid_search_attributes(log_reg_grid_results['model'])
```

```
-----
|      Best Estimator      |
-----
```

```
LogisticRegression(C=30, class_weight=None, dual=False, fit_intercept
=True,
                    intercept_scaling=1, max_iter=100, multi_class='warn',
                    n_jobs=None, penalty='l2', random_state=None, solver='warn',
                    tol=0.0001, verbose=0, warm_start=False)
```

```
-----
|    Best parameters      |
-----
```

```
Parameters of best estimator :

{'C': 30, 'penalty': 'l2'}
```

```
-----
| No of CrossValidation sets |
-----
```

```
Total numbre of cross validation sets: 3
```

```
-----
|      Best Score         |
-----
```

```
Average Cross Validate scores of best estimator :

0.9461371055495104
```

2. Linear SVC with GridSearch

```
In [15]: from sklearn.svm import LinearSVC
```

```
In [16]: parameters = {'C':[0.125, 0.5, 1, 2, 8, 16]}
lr_svc = LinearSVC(tol=0.00005)
lr_svc_grid = GridSearchCV(lr_svc, param_grid=parameters, n_jobs=-1, verbose=1
)
lr_svc_grid_results = perform_model(lr_svc_grid, X_train, y_train, X_test, y_t
est, class_labels=labels)
```

training the model..

```
C:\Users\Raftaar Singh\Anaconda3\lib\site-packages\sklearn\model_selection\_split.py:2053: FutureWarning: You should specify a value for 'cv' instead of relying on the default value. The default value will change from 3 to 5 in version 0.22.
```

```
warnings.warn(CV_WARNING, FutureWarning)
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
```

Fitting 3 folds for each of 6 candidates, totalling 18 fits

```
[Parallel(n_jobs=-1)]: Done 18 out of 18 | elapsed: 33.7s finished
```

```
C:\Users\Raftaar Singh\Anaconda3\lib\site-packages\sklearn\svm\base.py:931: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

Done

training_time(HH:MM:SS.ms) - 0:00:41.770372

Predicting test data

Done


testing time(HH:MM:SS:ms) - 0:00:00.010972

```
-----
| Accuracy |
-----
```

0.9657278588394977

```
-----
| Confusion Matrix |
-----
```

```
[[537  0  0  0  0  0]
 [ 2426 58  0  0  5]
 [  0 11520  1  0  0]
 [  0  0  0496  0  0]
 [  0  0  0  2413  5]
 [  0  0  0  16  1454]]
```



Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.97	0.87	0.92	491
STANDING	0.90	0.98	0.94	532
WALKING	0.96	1.00	0.98	496
WALKING_DOWNSTAIRS	1.00	0.98	0.99	420
WALKING_UPSTAIRS	0.98	0.96	0.97	471
micro avg	0.97	0.97	0.97	2947
macro avg	0.97	0.97	0.97	2947
weighted avg	0.97	0.97	0.97	2947

```
In [17]: print_grid_search_attributes(lr_svc_grid_results['model'])
```

```
-----
|      Best Estimator      |
|-----|
```

```
LinearSVC(C=2, class_weight=None, dual=True, fit_intercept=True,
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=5e-05,
verbose=0)
```

```
-----
|    Best parameters      |
|-----|
```

```
Parameters of best estimator :
```

```
{'C': 2}
```

```
-----
| No of CrossValidation sets |
|-----|
```

```
Total number of cross validation sets: 3
```

```
-----
|      Best Score         |
|-----|
```

```
Average Cross Validate scores of best estimator :
```

```
0.9469532100108814
```

3. Kernel SVM with GridSearch

```
In [18]: from sklearn.svm import SVC
parameters = {'C':[2,8,16],\
              'gamma': [ 0.0078125, 0.125, 2]}
rbf_svm = SVC(kernel='rbf')
rbf_svm_grid = GridSearchCV(rbf_svm,param_grid=parameters, n_jobs=-1)
rbf_svm_grid_results = perform_model(rbf_svm_grid, X_train, y_train, X_test, y_test, class_labels=labels)
```

training the model..

```
C:\Users\Raftaar Singh\Anaconda3\lib\site-packages\sklearn\model_selection\_split.py:2053: FutureWarning: You should specify a value for 'cv' instead of relying on the default value. The default value will change from 3 to 5 in version 0.22.
```

```
warnings.warn(CV_WARNING, FutureWarning)
```

Done

training_time(HH:MM:SS.ms) - 0:07:41.717459

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:04.439175

```
-----
| Accuracy |
-----
```

0.9626739056667798

```
-----
| Confusion Matrix |
-----
```

```
[[537  0  0  0  0  0]
 [ 0 441 48  0  0  2]
 [ 0 12 520  0  0  0]
 [ 0  0  0 489  2  5]
 [ 0  0  0  4 397 19]
 [ 0  0  0 17  1 453]]
```

Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.97	0.90	0.93	491
STANDING	0.92	0.98	0.95	532
WALKING	0.96	0.99	0.97	496
WALKING_DOWNSTAIRS	0.99	0.95	0.97	420
WALKING_UPSTAIRS	0.95	0.96	0.95	471
micro avg	0.96	0.96	0.96	2947
macro avg	0.96	0.96	0.96	2947
weighted avg	0.96	0.96	0.96	2947

```
In [19]: print_grid_search_attributes(rbf_svm_grid_results['model'])
```

```
-----
|      Best Estimator      |
|-----|
```

```
SVC(C=16, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=0.0078125, kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

```
-----
|    Best parameters      |
|-----|
```

```
Parameters of best estimator :
```

```
{'C': 16, 'gamma': 0.0078125}
```

```
-----
| No of CrossValidation sets |
|-----|
```

```
Total nombre of cross validation sets: 3
```

```
-----
|      Best Score        |
|-----|
```

```
Average Cross Validate scores of best estimator :
```

```
0.9440968443960827
```

4. Decision Trees with GridSearchCV

```
In [20]: from sklearn.tree import DecisionTreeClassifier
parameters = {'max_depth':np.arange(3,10,2)}
dt = DecisionTreeClassifier()
dt_grid = GridSearchCV(dt,param_grid=parameters, n_jobs=-1)
dt_grid_results = perform_model(dt_grid, X_train, y_train, X_test, y_test, class_labels=labels)
print_grid_search_attributes(dt_grid_results['model'])
```

training the model..

```
C:\Users\Raftaar Singh\Anaconda3\lib\site-packages\sklearn\model_selection\_split.py:2053: FutureWarning: You should specify a value for 'cv' instead of relying on the default value. The default value will change from 3 to 5 in version 0.22.
```

```
warnings.warn(CV_WARNING, FutureWarning)
```

Done

training_time(HH:MM:SS.ms) - 0:00:17.129530

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:00.009973

```
-----
| Accuracy |
-----
```

0.8646080760095012

```
-----
| Confusion Matrix |
-----
```

```
[[537  0  0  0  0  0]
 [ 0 386 105  0  0  0]
 [ 0  93 439  0  0  0]
 [ 0  0  0 471 17  8]
 [ 0  0  0 13 346 61]
 [ 0  0  0 73 29 369]]
```



Classification Report				

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.81	0.79	0.80	491
STANDING	0.81	0.83	0.82	532
WALKING	0.85	0.95	0.89	496
WALKING_DOWNSTAIRS	0.88	0.82	0.85	420
WALKING_UPSTAIRS	0.84	0.78	0.81	471
micro avg	0.86	0.86	0.86	2947
macro avg	0.86	0.86	0.86	2947
weighted avg	0.86	0.86	0.86	2947

```

-----
| Best Estimator |
-----

DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth
=7,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                        splitter='best')

```

```

-----
| Best parameters |
-----

Parameters of best estimator :

{'max_depth': 7}

```

```

-----
| No of CrossValidation sets |
-----

Total nombre of cross validation sets: 3

```

```

-----
| Best Score |
-----

Average Cross Validate scores of best estimator :

0.8375952121871599

```

5. Random Forest Classifier with GridSearch

```
In [21]: from sklearn.ensemble import RandomForestClassifier
params = {'n_estimators': np.arange(10,201,20), 'max_depth':np.arange(3,15,2)}
rfc = RandomForestClassifier()
rfc_grid = GridSearchCV(rfc, param_grid=params, n_jobs=-1)
rfc_grid_results = perform_model(rfc_grid, X_train, y_train, X_test, y_test, c
lass_labels=labels)
print_grid_search_attributes(rfc_grid_results['model'])
```

training the model..

```
C:\Users\Raftaar Singh\Anaconda3\lib\site-packages\sklearn\model_selection\_split.py:2053: FutureWarning: You should specify a value for 'cv' instead of relying on the default value. The default value will change from 3 to 5 in version 0.22.
```

```
warnings.warn(CV_WARNING, FutureWarning)
```

Done

training_time(HH:MM:SS.ms) - 0:09:22.684772

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:00.053856

```
-----
| Accuracy |
-----
```

0.9110960298608755

```
-----
| Confusion Matrix |
-----
```

```
[[537  0  0  0  0  0]
 [ 0 421 70  0  0  0]
 [ 0 47 485  0  0  0]
 [ 0  0  0 485  9  2]
 [ 0  0  0 33 341 46]
 [ 0  0  0 49  6 416]]
```

```
-----
| Classification Report |
-----
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.90	0.86	0.88	491
STANDING	0.87	0.91	0.89	532
WALKING	0.86	0.98	0.91	496
WALKING_DOWNSTAIRS	0.96	0.81	0.88	420
WALKING_UPSTAIRS	0.90	0.88	0.89	471
micro avg	0.91	0.91	0.91	2947
macro avg	0.91	0.91	0.91	2947
weighted avg	0.91	0.91	0.91	2947

```
-----
| Best Estimator |
-----
```

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion
='gini',
                        max_depth=7, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=70, n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

```
-----
| Best parameters |
-----
```

Parameters of best estimator :

```
{'max_depth': 7, 'n_estimators': 70}
```

```
-----
| No of CrossValidation sets |
-----
```

Total number of cross validation sets: 3

```
-----
| Best Score |
-----
```

Average Cross Validate scores of best estimator :

```
0.9148531011969532
```

6. Gradient Boosted Decision Trees With GridSearch

```
In [22]: from sklearn.ensemble import GradientBoostingClassifier
param_grid = {'max_depth': np.arange(5,8,1), \
              'n_estimators': np.arange(130,170,10)}
gbdt = GradientBoostingClassifier()
gbdt_grid = GridSearchCV(gbdt, param_grid=param_grid, n_jobs=-1)
gbdt_grid_results = perform_model(gbdt_grid, X_train, y_train, X_test, y_test,
class_labels=labels)
print_grid_search_attributes(gbdt_grid_results['model'])
```

training the model..

```
C:\Users\Raftaar Singh\Anaconda3\lib\site-packages\sklearn\model_selection\_split.py:2053: FutureWarning: You should specify a value for 'cv' instead of relying on the default value. The default value will change from 3 to 5 in version 0.22.
```

```
warnings.warn(CV_WARNING, FutureWarning)
```

Done

training_time(HH:MM:SS.ms) - 2:41:53.788148

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:00.198170

```
-----
| Accuracy |
-----
```

0.9222938581608415

```
-----
| Confusion Matrix |
-----
```

```
[[537  0  0  0  0  0]
 [ 0 398 91  0  0  2]
 [ 0 38 494  0  0  0]
 [ 0  0  0 483  7  6]
 [ 0  0  0 11 374 35]
 [ 0  1  0 32  6 432]]
```



Classification Report				

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.91	0.81	0.86	491
STANDING	0.84	0.93	0.88	532
WALKING	0.92	0.97	0.95	496
WALKING_DOWNSTAIRS	0.97	0.89	0.93	420
WALKING_UPSTAIRS	0.91	0.92	0.91	471
micro avg	0.92	0.92	0.92	2947
macro avg	0.92	0.92	0.92	2947
weighted avg	0.92	0.92	0.92	2947

```

-----
| Best Estimator |
-----

GradientBoostingClassifier(criterion='friedman_mse', init=None,
    learning_rate=0.1, loss='deviance', max_depth=5,
    max_features=None, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=160,
    n_iter_no_change=None, presort='auto', random_state=None,
    subsample=1.0, tol=0.0001, validation_fraction=0.1,
    verbose=0, warm_start=False)

```

```

-----
| Best parameters |
-----

Parameters of best estimator :

{'max_depth': 5, 'n_estimators': 160}

```

```

-----
| No of CrossValidation sets |
-----

Total nombre of cross validation sets: 3

```

```

-----
| Best Score |
-----

Average Cross Validate scores of best estimator :

0.9030195865070729

```

7. Comparing all models

```
In [23]: print('\n                Accuracy      Error')
print('                -----      -----')
print('Logistic Regression : {:.04}%      {:.04}%'.format(log_reg_grid_results['accuracy'] * 100, \
                                                         100-(log_reg_grid_results['accuracy'] * 100)))

print('Linear SVC          : {:.04}%      {:.04}% '.format(lr_svc_grid_results['accuracy'] * 100, \
                                                         100-(lr_svc_grid_results['accuracy'] * 100)))

print('rbf SVM classifier  : {:.04}%      {:.04}% '.format(rbf_svm_grid_results['accuracy'] * 100, \
                                                         100-(rbf_svm_grid_results['accuracy'] * 100)))

print('DecisionTree       : {:.04}%      {:.04}% '.format(dt_grid_results['accuracy'] * 100, \
                                                         100-(dt_grid_results['accuracy'] * 100)))

print('Random Forest      : {:.04}%      {:.04}% '.format(rfc_grid_results['accuracy'] * 100, \
                                                         100-(rfc_grid_results['accuracy'] * 100)))

print('GradientBoosting DT : {:.04}%      {:.04}% '.format(rfc_grid_results['accuracy'] * 100, \
                                                         100-(rfc_grid_results['accuracy'] * 100)))
```

	Accuracy -----	Error -----
Logistic Regression :	96.27%	3.733%
Linear SVC :	96.57%	3.427%
rbf SVM classifier :	96.27%	3.733%
DecisionTree :	86.46%	13.54%
Random Forest :	91.11%	8.89%
GradientBoosting DT :	91.11%	8.89%

We can choose **Logistic regression** or **Linear SVC** or **rbf SVM**.

Conclusion :

In the real world, domain-knowledge, EDA and feature-engineering matter most.

